

POSC 3XX: AI Governance & Policy

Fall 2026

Faculty Information

Instructor: David P. Adams, Ph.D.

Office: 516 Gordon Hall

Phone/Text: (657) 278-4770

Email: dpadams@fullerton.edu

Website: dadams.io

Office Hours: Mondays and Wednesdays, [TBD], and by [appointment](#).

Schedule meetings: dadams.io/appointments

Course Communication

All course announcements and communications will be sent via *Canvas* and university email. Students are responsible for regularly checking their *Canvas* notifications and email, and for ensuring that Canvas notifications are set to receive course messages. Students are expected to check *Canvas* and their email at least once daily.

Response time: I will strive to respond to all student emails and *Canvas* messages within 24 hours, except on weekends and holidays. If you have not received a response within 24 hours, please send a follow-up message. If you are still waiting after 48 hours, contact me via phone or SMS at (657) 278-4770.

Technical Problems

If you encounter any technical difficulties, contact the instructor immediately to document the problem. Then contact: [student IT help desk](#), [email](#), phone (657) 278-8888, walk-in [student genius center](#), or online chat via the [portal](#) (“Online IT Help” then “Live Chat”).

For issues with Canvas: Canvas Support Hotline = (657) 278-8888, [search the CSUF Canvas Guides](#), or [report a problem](#).

Alternative submission: If you cannot submit an assignment via *Canvas*, contact the professor as soon as possible to document the issue and arrange an alternative.

Course Information

Prefix, number, title: POSC 3XX, *AI Governance & Policy*

Meeting times: In-Person, Monday & Wednesday, [Time TBD], [Room TBD]

Units: 3 **Schedule Code:** [TBD]

Course requisite(s): POSC 100 (American Government) or equivalent; completion of GE Area 4A (required as a prerequisite for all Area 4U courses)

GE designation: Area 4U, Explorations in Social and Behavioral Sciences [pending certification]

Cross-listing: POSC/CPSC 3XX [pending approval]

Policy regarding the use of generative AI: See the *Policy on the Use of Generative AI and Other Technology* section below.

Course materials and equipment: Canvas; access to course readings; laptop recommended

Required texts: Eubanks, *Automating Inequality* (see *Required Texts* below); article packet provided through Canvas and the Pollak Library

Catalog Description

Governance and politics of artificial intelligence. Regulatory institutions, public-sector AI use, procurement, accountability, liability, safety and alignment, infrastructure, cybersecurity, democracy and comparative regulation. No programming required.

Course Description

Artificial intelligence is often discussed as a technical problem or an ethical dilemma. This course treats it as something else: a governance problem. Who regulates AI, at which level of government, with what authority, and with what actual capacity to enforce? What

happens when public agencies (which mostly buy AI systems rather than build them) delegate consequential decisions about benefits, policing, hiring, and health care to automated systems? And who is responsible when those systems fail people?

The course examines the institutions, laws, and administrative processes through which societies attempt to govern AI. Students will study the fragmented American regulatory landscape alongside the European Union's omnibus approach; algorithmic decision-making inside government agencies; the physical footprint of AI infrastructure (data centers, energy, water, land use) as a live federalism problem; liability regimes for AI-caused harm; AI as a cybersecurity concern; the distinctive governance challenges of increasingly autonomous (agentic) and embodied (physical) AI systems—where *alignment*, the gap between the objectives a system is built to optimize and the goals and values its human principals actually intend, becomes a governance problem as agents gain the autonomy to act on their own; and the emerging effects of generative AI on democratic institutions themselves.

Rather than building AI systems, students will learn to *audit* them. The semester-long audit project asks each student to evaluate a real AI system: its documentation, its governance, its failure modes, and the accountability structures (or gaps) surrounding it. This is the skill set behind emerging roles in AI governance, compliance, and oversight, and it requires no programming whatsoever.

The course is designed for students in computer science, public administration, political science, and related fields. Technical and institutional expertise are treated as complementary: students with technical backgrounds will learn to connect system design and documentation to law, public administration, procurement, and accountability, while students from policy and social science backgrounds will build the AI literacy needed to evaluate technical claims.

The central question guiding the course is:

When an algorithm makes a decision that harms someone, who is accountable—and what would it take for that answer to be satisfying?

Student Learning Outcomes

By the end of the semester, students will be able to:

1. Explain how regulatory authority over AI is distributed across levels and branches of American government.
2. Compare sectoral (U.S.) and omnibus (EU) approaches to AI regulation and evaluate their institutional trade-offs.
3. Analyze the use of algorithmic decision-making systems within public agencies and their effects on administrative discretion and accountability.
4. Evaluate liability and accountability regimes for harms caused by automated systems.
5. Assess the governance challenges posed by AI infrastructure, including data center siting and resource demands.
6. Analyze the effects of generative AI on democratic processes and institutions.
7. Evaluate emerging governance challenges posed by increasingly autonomous and embodied AI systems, including alignment, human oversight, and safety assurance.
8. Translate technical documentation about AI systems into policy-relevant claims about risk, accountability, oversight, and institutional responsibility.
9. Conduct a structured audit of a real AI system's documentation, governance, and accountability mechanisms.

General Education Information

Requirements Satisfied

This course satisfies General Education Area 4U, Explorations in Social and Behavioral Sciences. Because courses in Area 4U build upon the learning objectives in Area 4A, completion of Area 4A is required as a prerequisite for this course. The writing assignments in this course, including the scaffolded AI system audit project described below, meet the requirement of UPS 411.201:

Writing assignments in General Education courses shall involve the organization and expression of complex data or ideas and careful and timely evaluations of writing so that deficiencies are identified, and suggestions for improvement and/or for means of remediation are offered. Evaluations of the student's writing competence shall determine the final course grade....

A grade of D (1.0) or higher is required to meet this General Education requirement. A grade of D- (0.7) or below will not satisfy this General Education Requirement.

General Education Student Learning Goals

Students completing courses in Area 4U shall:

1. Examine problems, issues, and themes in the social sciences in greater depth; in a variety of cultural, historical, and geographical contexts; and from different disciplinary and interdisciplinary perspectives.
2. Analyze and critically evaluate the application of social science concepts and theories to particular historical, contemporary, and future problems or themes, such as economic and environmental sustainability, globalization, poverty, and social justice.
3. Analyze and critically evaluate constructs of cultural differentiation, including ethnicity, gender, race, class, and sexual orientation, and their effects on the individual and society.
4. Apply theories and concepts from the social sciences to address historical, contemporary, and future problems confronting communities at different geographical scales, from local to global.

Required Texts

Required Book

- Eubanks, Virginia. 2018. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. New York, NY: St. Martin's Press.

Article Packet

All other readings (journal articles, investigative journalism, regulatory documents, and case materials) are provided through Canvas and the Pollak Library at no cost. The article packet carries developments in AI policy since the book's publication and is updated each offering of the course. Canvas readings are required and carry equal weight to assigned book chapters.

Student Resources Website

It is the student's responsibility to read and understand the required and important [student information for course syllabi](#). Included is information about:

- University learning goals and General Education learning objectives
- Students' rights to accommodations
- Campus student support resources and academic integrity
- Emergency preparedness; library and IT services
- Software privacy, accessibility statement, diversity statement, and land acknowledgement
- Final exam schedule and semester calendar

Course Requirements

Course Format

This is an in-person lecture and discussion course. Students complete assigned readings before class; class time is divided between lecture, structured discussion, and case analysis. No programming, mathematics, or prior technical coursework is required or assumed. Students who bring technical experience are encouraged to use it, but the central work of the course is governance analysis: connecting AI systems to institutions, law, procurement, oversight, and democratic accountability.

Graded Work

- **AI System Audit Project (60% total):** The semester's central assignment. Each student selects a real, documented AI system (from an instructor-provided list or by petition) and conducts a structured governance audit: what the system does, who deploys it, what its documented failure modes are, which laws and oversight mechanisms apply, and where the accountability gaps sit. Students may audit systems in public benefits, policing and criminal justice, health care, hiring, education, infrastructure, cybersecurity, platform governance, elections and democracy, or other instructor-approved domains. The project unfolds in stages:

- **Audit Checkpoint (10%):** System selected, documentation gathered, audit framework outlined. Due Monday, October 19, 11:59pm on Canvas.
 - **Audit Progress Memo (5%):** A two-page memo reporting documentation reviewed, preliminary findings, and open questions, with instructor feedback returned before the audit fair. Due Wednesday, November 18.
 - **Audit Brief and Fair (15%):** Students prepare a concise audit brief and discuss their findings in a structured in-class audit fair on Monday, November 30 and Wednesday, December 2. Students also review classmates' audits (peer-response assignments pair students across disciplinary backgrounds where possible) and use the feedback to revise the final report.
 - **Final Audit Report (30%):** Written report due during finals week (December 14–18); exact date and time assigned by Registrar. The report must include a section applying at least two frameworks from the second half of the course—liability, security, autonomous-agent oversight, or democratic impact—to the audited system.
- **Midterm Examination (20%):** In-class examination on Wednesday, October 21, covering Parts I–III.
 - **Participation and Weekly Responses (20% total):** Measured in two parts:
 - **Weekly Reading Responses (12%):** A short response to the week's readings (150–250 words), posted to the Canvas discussion board by 11:59pm the Sunday before Monday's class. Twelve responses are assigned across the semester; the best ten count. Graded credit/no-credit against a posted rubric (engages the actual readings, makes a claim, asks a question).
 - **Attendance and In-Class Engagement (8%):** Attendance is recorded each session. Engagement means arriving prepared and contributing to discussion, case analysis, and the audit fair. Up to two absences are excused without documentation; participation in in-class activities cannot be made up (see attendance policy below).

Grading Policies and Standards

a. **Grading scale:** See Table 1 for the full letter-grade percentage scale used in this course.

Table 1: Grade scale

Grade	Percent	Grade	Percent
A+	98.0–100.0	C+	77.0–79.9
A	93.0–97.9	C	73.0–76.9
A-	90.0–92.9	C-	70.0–72.9
B+	87.0–89.9	D+	67.0–69.9
B	83.0–86.9	D	63.0–66.9
B-	80.0–82.9	D-	60.0–62.9
		F	0.0–59.9

b. Required course assignments: See Table 2 for assignment weights and due dates.

Table 2: Assignment weighting

Assignment	Weight	Due
Audit Checkpoint	10%	Monday, October 19
Midterm Examination	20%	Wednesday, October 21
Audit Progress Memo	5%	Wednesday, November 18
Audit Brief and Fair	15%	November 30 and December 2
Final Audit Report	30%	Finals Week, December 14–18
Participation and Weekly Responses	20%	Ongoing
Total	100%	

c. Attendance and participation policy: Students are expected to attend all in-person sessions. If you are unable to attend, notify the professor in advance. You are responsible for obtaining any materials or information covered during absences. Participation in in-class activities cannot be made up.

d. Examination dates:

- Midterm Examination: Wednesday, October 21 (in-class)
- Final Audit Report due during finals week; exact date and time assigned by Registrar (December 14–18)

e. Late work: Late submissions are accepted up to 72 hours after the deadline with a 10% per-day penalty, except the audit fair activities and final report, which cannot be submitted

late absent documented emergency. Contact the professor *before* a deadline whenever possible.

f. Authentication of student work: Students may be required to submit their work to a plagiarism detection service. Cal State Fullerton uses Turnitin®. Students should be aware that submitted work may be checked for authenticity and originality.

g. Extra credit: There are no extra credit assignments in this course.

h. Retention of student work: Work submitted for a grade, whether as a hard copy or through Canvas, shall be retained for a reasonable time after the semester ends, not to exceed the last day of the subsequent semester. Students have the right to review graded work in the presence of the instructor. (UPS 320.005)

Academic Integrity

Students are expected to adhere to the highest standards of academic integrity. Any student found to have engaged in academic dishonesty will be subject to the sanctions described in the [Academic Dishonesty Policy](#) (UPS 300.021). Academic dishonesty includes, but is not limited to, cheating, plagiarism, fabrication, facilitating academic dishonesty, and submitting previously graded work without prior authorization. Students are expected to be familiar with the university's policy and to adhere to it in all aspects of this course.

Policy on the Use of Generative AI and Other Technology

Generative AI (including large language models, image generators, and other tools) is permitted in this course, but use must be transparent, intentional, and in service of learning. The core principle is simple: **you must do the intellectual work of this course**. AI can amplify your thinking, but not replace it. (Yes, the irony of an AI-use policy in an AI policy course is noted. Consider it your first case study.)

You may use general-purpose AI tools (ChatGPT, Claude, Gemini, etc.) subject to the rules below. The permitted/not-permitted list applies equally regardless of which tool you use.

Permitted uses:

- Brainstorming and outlining arguments
- Explaining concepts you don't understand (then explaining it back in your own words)
- Literature searching and summarizing sources
- Editing, proofreading, and revising your work
- Sanity-checking your analysis or logic
- Generating synthetic examples or test cases for your ideas

Not permitted:

- Using AI to generate your analysis, arguments, or conclusions
- Submitting AI-generated text as your own writing
- Using AI to avoid engaging with course concepts or readings
- Letting AI do the intellectual heavy lifting (interpreting sources, building arguments, synthesizing ideas)

Disclosure requirement:

If you use AI tools in ways beyond basic editing, you must disclose your use. Include a brief note at the end of your assignment explaining what tools you used and how (e.g., “I used Claude to help organize my outline and check the logic of my argument in Section 3”). This is not a confession: it’s transparency about your process.

What this means:

The goal of this course is for *you* to learn to think like a policy analyst about AI governance and to develop your own informed arguments. AI is a tool that can enhance that learning if used thoughtfully. Using it to avoid thinking will undermine your own education and violates academic integrity. Questions about what constitutes appropriate use? Ask before you submit.

Technical Competencies

Students need:

- Proficiency with Canvas, including submitting assignments and accessing course materials
- Ability to use university email and Canvas messages for course communication
- Basic word processing skills and ability to export documents to PDF

Calendar of Topics / Schedule of Classes

We will follow the schedule below as closely as possible. If adjustments are needed, you will receive advance notice in class and on *Canvas*.

Important dates:

- Audit Checkpoint due: Monday, October 19 (11:59pm on Canvas)
- Midterm Examination: Wednesday, October 21 (in-class)
- Audit Progress Memo due: Wednesday, November 18
- Audit Fair: Monday, November 30 and Wednesday, December 2
- Final Audit Report: finals week, December 14–18 (Registrar-assigned)
- Labor Day: Monday, September 7 (no class)
- Veterans Day: Wednesday, November 11 (no class)
- Thanksgiving Break: November 23–27 (no class)

PART I: Foundations

8/24 – Week 1: Why is this a Governance Problem?

Monday 8/24: The Governance Frame

- AI as a governance problem, not (only) a technical or ethical one
- The central question: when an algorithm harms someone, who is accountable?
- Course framing: audit, don't build

Wednesday 8/26: Institutions, Enforcement, Capacity

- Who regulates, with what authority, and with what capacity to enforce
- Institutional literacy refresher: agencies, rulemaking, courts, and federalism—the policy-side counterpart to Week 2's AI literacy
- Why governance chronically lags technology—and whether it must

Readings

- Bullock (2019), "Artificial Intelligence, Discretion, and Bureaucracy," *American Review of Public Administration*
- Wheeler (2023), "The Three Challenges of AI Regulation," Brookings Institution

8/31 – Week 2: AI Literacy for Policy and Governance

Monday 8/31: From Rules to Learning

- What machine learning systems do (and don't do): from symbolic AI to machine learning to deep learning
- Training data, model behavior, and failure modes in plain language
- The black-box problem, and why "the algorithm decided" is never a complete sentence

Wednesday 9/2: The Frontier—Generative, Agentic, and Physical AI

- Large language models and generative AI: capabilities, limits, and commercial maturity
- Agentic AI: systems that plan, use tools, and pursue goals with autonomy
- Physical AI: embodied systems that sense and act on the real world
- Why each generation of AI produces a new generation of policy problems

Readings

- Narayanan and Kapoor (2024), *AI Snake Oil*, introduction (excerpt)
- Mitchell et al. (2019), “Model Cards for Model Reporting,” plus one current production model card (skim)

Guest Lecture

- Department of Computer Science faculty [TBD]

PART II: Governance and Federalism

9/7 – Week 3: Who Regulates? The U.S. Sectoral Patchwork

Monday September 7 is Labor Day — no class. Wednesday September 9 only.

Topics

- Distributed regulatory authority: agencies, states, courts
- Federalism and jurisdictional overlap
- Why enforcement capacity lags statutory ambition

Readings

- Engstrom, Ho, Sharkey, and Cuéllar (2020), *Government by Algorithm* (ACUS report), executive summary and one agency case study
- Colorado AI Act (SB24-205) explainer; NCSL AI legislation tracker (skim)

9/14 – Week 4: Governance Made Concrete: AI Infrastructure

Monday 9/14: The Physical Footprint of AI

- Data center siting, energy, water, and land use
- Environmental impacts of computational resource consumption

Wednesday 9/16: Federalism with Teeth

- Local zoning authority versus hyperscale capital
- Who decides where AI physically lives

Readings

- Current investigative reporting on data center energy and water demands (updated each offering)
- Hogan (2015), “Data Flows and Water Woes: The Utah Data Center,” *Big Data & Society*
- Case dossier: local coverage of a current data center siting dispute

9/21 – Week 5: The EU AI Act as Counterfactual

Monday 9/21: The EU AI Act

- Omnibus versus sectoral regulatory design
- Risk-tier logic and conformity assessment

Wednesday 9/23: Regulatory Diffusion and the Three-Way Comparison

- The “Brussels effect” and regulatory diffusion
- Comparative snapshot: EU risk-based regulation, U.S. sectoral patchwork, China’s application-specific rules

Readings

- Veale and Zuiderveen Borgesius (2021), “Demystifying the Draft EU Artificial Intelligence Act” (excerpt)
- Bradford (2023), *Digital Empires*, three-empires framework chapter (excerpt)

Assignments

- **Audit Project launch:** system list distributed; selection guidance in class

PART III: The State as AI User

9/28 – Week 6: Algorithmic Decision-Making Inside Agencies

Monday 9/28: Street-Level Bureaucracy Meets Automated Eligibility

- How algorithmic systems enter public agencies
- Automated eligibility and the transformation of frontline work

Wednesday 9/30: Algorithmic Bias, Fairness, and Who Gets Flagged

- Algorithmic bias and discrimination: how race, class, gender, and disability shape who gets flagged
- Administrative discretion, legibility, and disparate impact in automated decisions
- What automation does to accountability inside agencies

Readings

- Eubanks, introduction and chapters 1–2
- Bovens and Zouridis (2002), “From Street-Level to System-Level Bureaucracies,” *Public Administration Review*

10/5 – Week 7: Case Deep-Dive: When Automated Systems Fail People

Monday 10/5: Anatomy of an Algorithmic Failure

- Design, deployment, denial, discovery

Wednesday 10/7: The Case in Depth

- Case: [Dutch childcare benefits scandal *or* COMPAS risk assessment—one, in depth]
- Fairness contested: competing definitions of a “fair” algorithm and who bears the burden of error

Readings

- Eubanks, chapters 3–4
- Case dossier (one assigned per offering): COMPAS—Angwin et al. (2016), “Machine Bias,” *ProPublica*, plus the Northpointe rebuttal exchange; *or* Dutch childcare benefits scandal—Amnesty International (2021), *Xenophobic Machines* (excerpt)

10/12 – Week 8: Procurement and the Governance Toolkit

Monday 10/12: Procurement: Where Governance Really Happens

- Agencies buy AI; they rarely build it
- Contracting, vendor opacity, and trade-secret claims against transparency
- Procurement as the real site of AI governance inside government

Wednesday 10/14: Governance Mechanisms in Practice

- The emerging AI assurance toolkit: risk assessments, model evaluation, red-team testing, human-in-the-loop oversight, continuous monitoring and auditing
- Governance across the AI lifecycle: design, development, deployment, monitoring, retirement
- How these mechanisms map onto your audit framework

Readings

- Eubanks, chapter 5 and conclusion
- Pahlka (2023), *Recoding America* (one chapter)
- NIST (2023), *AI Risk Management Framework* (AI RMF 1.0), core functions (skim)

10/19 – Week 9: Review and Midterm

Monday 10/19: Midterm Review and Audit Clinic

- Structured review of Parts I–III
- Open audit clinic: final framework troubleshooting before tonight's checkpoint deadline

Assignments

- **Monday October 19: Audit Checkpoint due** (11:59pm on Canvas; system selected, documentation gathered, audit framework outlined)
- **Wednesday October 21: Midterm Examination** (in-class, covering Parts I–III)

PART IV: Harm, Liability, and Security

10/26 – Week 10: When the Algorithm Hurts Someone

Monday 10/26: Liability for Automated Decisions

- Health care denial-of-care algorithms; hiring discrimination suits
- Tort law meets administrative oversight: who pays, under which theory

Wednesday 10/28: Physical AI—When Software Gets a Body

- Autonomous vehicles, drones, and robots: liability when AI causes physical harm
- Safety assurance, testing and certification, and product liability for embodied systems
- Autonomous vehicle litigation and theories of liability

Readings

- Obermeyer et al. (2019), “Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations,” *Science* (excerpt through main results; methods appendix optional)
- Ross and Herman (2023), “Denied by AI,” *STAT News* investigation
- Case dossier: current autonomous-vehicle litigation coverage with a short liability explainer (updated each offering)

11/2 – Week 11: Autonomy and Security

Monday 11/2: Governing Autonomous Agents

- The oldest problem in political science meets the newest technology: AI agents as a principal–agent problem
- The alignment problem: systems optimize a specified mathematical objective, but a stated objective rarely captures the full range of human goals and values—so the agent can satisfy the letter of the goal in unintended or harmful ways, and the more autonomy it has to plan and act, the higher the stakes
- Translating the technical vocabulary: alignment as goal fidelity, oversight as police patrols versus fire alarms, corrigibility as the power to correct or recall

- Who is accountable when an agent acts: agency law, early chatbot liability rulings, and delegation of authority to software
- Legislating the principal–agent relationship: human-oversight mandates (EU AI Act, Article 14)

Wednesday 11/4: AI as Attack Surface

- Adversarial inputs and security of automated pipelines
- AI in cyber offense and defense; autonomous systems and national security
- Governing systems that can be manipulated

Readings

- Christian (2020), *The Alignment Problem*, prologue and chapter 1 (excerpt)
- Current reporting on AI agent accountability and liability (updated each offering)
- CISA/NCSC (2023), *Guidelines for Secure AI System Development* (skim)

Guest Lecture

- Cybersecurity faculty, Department of Computer Science [TBD] (Wednesday)

PART V: Democracy and Development

11/9 – Week 12: AI and Democratic Institutions

Monday November 9 only — Wednesday November 11 is Veterans Day; no class.

Topics

- Synthetic media and election regulation
- AI-generated astroturf in public comment processes
- Can notice-and-comment survive generative AI?

Readings

- New York State Office of the Attorney General (2021), *Fake Comments: How U.S. Companies & Partisans Hack Democracy to Undermine Your Voice*, executive summary
- Current explainer on AI and election integrity; NCSL deepfake election-law tracker (skim; updated each offering)

11/16 – Week 13: Responsible Development and the Future of Work

Monday 11/16: Responsible Development Across Sectors

- Health care, hiring, and content moderation compared
- Who monitors the monitors: regulatory capture and information asymmetry
- Comparative closer: China's algorithmic recommendation regulations

Wednesday 11/18: AI, Work, and Inequality

- Labor displacement and workforce transformation across white-collar and physical work
- Concentration of AI capability and winner-take-all dynamics
- What governments can (and cannot) do about it

Readings

- Stanford DigiChina (2022), translation and commentary: China's Internet Information Service Algorithmic Recommendation Management Provisions (excerpt)
- Brynjolfsson (2022), "The Turing Trap: The Promise & Peril of Human-Like Artificial Intelligence," *Daedalus*

Assignments

- **Wednesday November 18: Audit Progress Memo due** (two-page memo: documentation reviewed, preliminary findings, open questions; instructor feedback returned before the audit fair)

PART VI: Synthesis

Thanksgiving Break: November 23–27. No class.

11/30 – Week 14: Audit Fair

Assignments

- **Monday November 30: Audit Fair, Day 1:** Group A presents concise audit briefs; Group B circulates, asks questions, and completes peer responses
- **Wednesday December 2: Audit Fair, Day 2:** Group B presents concise audit briefs; Group A circulates, asks questions, and completes peer responses
- Peer-response assignments pair students across disciplinary backgrounds where possible: technical students review policy-side audits and vice versa

Notes

- Use instructor feedback on the Audit Progress Memo to sharpen your brief before presenting

12/7 – Week 15: Synthesis and Final Report Workshop

Monday 12/7: Synthesis

- Cross-case patterns from the audit fair: recurring failure modes, accountability gaps, and oversight tools
- Governing what doesn't exist yet: AGI debates, speculative risk, and international coordination
- What does a good AI governance regime look like a decade from now?
- Closing the loop on Week 1: institutions, enforcement, capacity

Assignments

- **Wednesday December 9: Final Audit Report workshop** (students use audit fair feedback to revise their final reports)

12/14 – Week 16: Finals Week

Deliverables

- **Final Audit Report due:** Date and time assigned by Registrar (December 14–18)

Appendix: Reading List

Full citations for the assigned readings listed in the course schedule above. All items other than the required book are provided through *Canvas* and the Pollak Library at no cost. Readings marked in the schedule as updated each offering (current investigative reporting, case dossiers, and regulatory trackers) are posted to *Canvas* at the start of the semester and are not listed here.

Amnesty International. 2021. Xenophobic Machines: Discrimination through Unregulated Use of Algorithms in the Dutch Childcare Benefits Scandal. Technical Report EUR 35/4686/2021 Amnesty International. Dossier B; excerpt assigned.

URL: <https://www.amnesty.org/en/documents/eur35/4686/2021/en/>

Angwin, Julia, Jeff Larson, Surya Mattu and Lauren Kirchner. 2016. “Machine Bias.” ProPublica. Dossier A; paired with the Northpointe technical rebuttal exchange.

URL: <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>

Bovens, Mark and Stavros Zouridis. 2002. “From Street-Level to System-Level Bureaucracies: How Information and Communication Technology Is Transforming Administrative Discretion and Constitutional Control.” *Public Administration Review* 62(2):174–184.

Bradford, Anu. 2023. *Digital Empires: The Global Battle to Regulate Technology*. New York, NY: Oxford University Press. Three-empires framework chapter assigned as excerpt.

- Brynjolfsson, Erik. 2022. "The Turing Trap: The Promise & Peril of Human-Like Artificial Intelligence." *Daedalus* 151(2):272–287.
- Bullock, Justin B. 2019. "Artificial Intelligence, Discretion, and Bureaucracy." *The American Review of Public Administration* 49(7):751–761.
- Christian, Brian. 2020. *The Alignment Problem: Machine Learning and Human Values*. New York, NY: W. W. Norton. Prologue and chapter 1 assigned as excerpt.
- Colorado General Assembly. 2024. "Consumer Protections for Artificial Intelligence (SB24-205)." Colorado Revised Statutes tit. 6, art. 1, pt. 17.
URL: <https://leg.colorado.gov/bills/sb24-205>
- Cybersecurity and Infrastructure Security Agency and National Cyber Security Centre. 2023. Guidelines for Secure AI System Development. Technical report CISA and NCSC. Assigned as skim.
URL: <https://www.cisa.gov/resources-tools/resources/guidelines-secure-ai-system-development>
- De Vynck, Gerrit. 2026. "One of Sci-Fi's Most Difficult Questions about AI Is Becoming Real." The Washington Post.
URL: <https://www.washingtonpost.com/technology/2026/07/13/one-sci-fis-most-difficult-questions-about-ai-is-becoming-real/>
- DigiChina. 2022. "Translation: Internet Information Service Algorithmic Recommendation Management Provisions—Effective March 1, 2022." Stanford University Cyber Policy Center.
URL: <https://digichina.stanford.edu/work/translation-internet-information-service-algorithmic-recommendation-management-provisions-effective-march-1-2022/>
- Engstrom, David Freeman, Daniel E. Ho, Catherine M. Sharkey and Mariano-Florentino Cuéllar. 2020. Government by Algorithm: Artificial Intelligence in Federal Administrative Agencies. Technical report Administrative Conference of the United States. Executive summary and one agency case study assigned.
URL: <https://www.acus.gov/report/government-algorithm-artificial-intelligence-federal-administrative-agencies>
- Eubanks, Virginia. 2018. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. New York, NY: St. Martin's Press. Required book; purchase or library reserve.

- Hogan, M  l. 2015. "Data Flows and Water Woes: The Utah Data Center." *Big Data & Society* 2(2).
- Mitchell, Margaret, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji and Timnit Gebru. 2019. Model Cards for Model Reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* '19)*. New York, NY: Association for Computing Machinery pp. 220–229.
- Narayanan, Arvind and Sayash Kapoor. 2024. *AI Snake Oil: What Artificial Intelligence Can Do, What It Can't, and How to Tell the Difference*. Princeton, NJ: Princeton University Press. Introduction assigned as excerpt.
- National Conference of State Legislatures. N.d. "Artificial Intelligence 2025 Legislation." NCSL.
URL: <https://www.ncsl.org/technology-and-communication/artificial-intelligence-2025-legislation>
- National Institute of Standards and Technology. 2023. Artificial Intelligence Risk Management Framework (AI RMF 1.0). Technical Report NIST AI 100-1 U.S. Department of Commerce. Core functions assigned as skim.
URL: <https://www.nist.gov/itl/ai-risk-management-framework>
- New York State Office of the Attorney General. 2021. Fake Comments: How U.S. Companies & Partisans Hack Democracy to Undermine Your Voice. Technical report Office of the New York State Attorney General. Executive summary assigned.
URL: <https://ag.ny.gov/sites/default/files/oag-fakecomments-report.pdf>
- Obermeyer, Ziad, Brian Powers, Christine Vogeli and Sendhil Mullainathan. 2019. "Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations." *Science* 366(6464):447–453. Assigned through main results; methods appendix optional.
- Pahlka, Jennifer. 2023. *Recoding America: Why Government Is Failing in the Digital Age and How We Can Do Better*. New York, NY: Metropolitan Books. One chapter assigned as excerpt.
- Ross, Casey and Bob Herman. 2023. "Denied by AI: How Medicare Advantage Plans Use Algorithms to Cut Off Care for Seniors in Need." STAT News. First installment of the four-part "Denied by AI" series; Pulitzer Prize finalist.

URL: <https://www.statnews.com/2023/03/13/medicare-advantage-plans-denial-artificial-intelligence/>

Veale, Michael and Frederik Zuiderveen Borgesius. 2021. "Demystifying the Draft EU Artificial Intelligence Act: Analysing the Good, the Bad, and the Unclear Elements of the Proposed Approach." *Computer Law Review International* 22(4):97–112.

Wheeler, Tom. 2023. "The Three Challenges of AI Regulation." Brookings Institution.

URL: <https://www.brookings.edu/articles/the-three-challenges-of-ai-regulation/>