

Reading Packet

POSC 316 / CPSC 316 – AI Governance and Policy, Fall 2028

David P. Adams, Ph.D.

Fall 2028

Cal State Fullerton.

How This Packet Works

This course assigns **one book** and an **article packet**. Every item other than the book is provided at no cost through *Canvas* and the Pollak Library.

Each week pairs **one scholarly reading with one accessible reading** — journalism, a plain-language explainer, or a primary regulatory document. That cap is deliberate: expect roughly **40–60 pages per week**. Weeks 6 through 8 lean on the book, so the article load drops.

There are **twelve reading weeks**, matching the twelve weekly reading responses. Weeks 9, 14, 15, and 16 assign no new reading so you can concentrate on the midterm, the audit fair, and your final report.

Some readings are marked *updated each offering*. AI policy moves quickly, so current investigative reporting, live case dossiers, and legislative trackers are refreshed each time the course runs and posted to *Canvas* at the start of the semester.

Short citations appear under each week. **Full citations for everything are listed at the end of this packet.**

Required Book

Eubanks, Virginia. 2018. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. New York: St. Martin's Press.

This is the only item you need to obtain yourself. Copies are on reserve at the Pollak Library.

Part I: Foundations

Week 1 (8/21) — Why Is This a Governance Problem?

Monday — The Governance Frame

- Bullock (2019), "Artificial Intelligence, Discretion, and Bureaucracy," *American Review of Public Administration* 49(7)

Wednesday — Institutions, Enforcement, Capacity

- Wheeler (2023), "The Three Challenges of AI Regulation," Brookings Institution

Week 2 (8/28) — AI Literacy for Policy and Governance

Monday — From Rules to Learning

- Narayanan and Kapoor (2024), *AI Snake Oil*, introduction (excerpt)

Wednesday — The Frontier: Generative, Agentic, and Physical AI

- Mitchell et al. (2019), "Model Cards for Model Reporting"
- One current production model card, skim (*updated each offering*)

Part II: Governance and Federalism

Week 3 (9/4) — Who Regulates? The U.S. Sectoral Patchwork

Labor Day Monday — Wednesday 9/6 only.

- Engstrom, Ho, Sharkey, and Cuéllar (2020), *Government by Algorithm* (ACUS report): executive summary and one agency case study
- Colorado AI Act (SB24-205) explainer; NCSL AI legislation tracker, skim (*updated each offering*)

Week 4 (9/11) — Governance Made Concrete: AI Infrastructure

Monday — The Physical Footprint of AI

- Nicoletti, Ma, and Bass (2025), “The AI Boom Is Draining Water from the Areas That Need It Most,” *Bloomberg*

Wednesday — Federalism with Teeth

- Hogan (2015), “Data Flows and Water Woes: The Utah Data Center,” *Big Data & Society* 2(2)
- **Case dossier:** Main (2025), “‘The City That Draws the Line’: One Arizona Community’s Fight against a Huge Datacenter,” *The Guardian*, with follow-on coverage (*updated each offering*)

Week 5 (9/18) — The EU AI Act as Counterfactual

Monday — The EU AI Act

- Veale and Zuiderveen Borgesius (2021), “Demystifying the Draft EU Artificial Intelligence Act,” *Computer Law Review International* 22(4) (excerpt: risk tiers and conformity assessment)

Wednesday — Regulatory Diffusion and the Three-Way Comparison

- Bradford (2023), *Digital Empires*, three-empires framework chapter (excerpt)

Part III: The State as AI User

Week 6 (9/25) — Algorithmic Decision-Making Inside Agencies

- **Eubanks**, introduction and chapters 1–2
- Bovens and Zouridis (2002), “From Street-Level to System-Level Bureaucracies,” *Public Administration Review* 62(2)

Week 7 (10/2) — Case Deep-Dive: When Automated Systems Fail People

- **Eubanks**, chapters 3–4
- **Case dossier**, one assigned per offering:
 - *COMPAS* — Angwin, Larson, Mattu, and Kirchner (2016), “Machine Bias,” *ProPublica*, with the Northpointe technical rebuttal exchange
 - *Dutch childcare benefits scandal* — Amnesty International (2021), *Xenophobic Machines* (excerpt)

Week 8 (10/9) — Procurement and the Governance Toolkit

- **Eubanks**, chapter 5 and conclusion
- Pahlka (2023), *Recoding America*, one chapter (excerpt)
- NIST (2023), *AI Risk Management Framework (AI RMF 1.0)*, core functions, skim

Week 9 (10/16) — Review and Midterm

No new reading. Audit Checkpoint due Monday; midterm examination Wednesday.

Part IV: Harm, Liability, and Security

Week 10 (10/23) — When the Algorithm Hurts Someone

Monday — Liability for Automated Decisions

- Obermeyer, Powers, Vogeli, and Mullainathan (2019), “Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations,” *Science* 366(6464) — read through the main results; the methods appendix is optional
- Ross and Herman (2023), “Denied by AI,” *STAT News*

Wednesday — Physical AI: When Software Gets a Body

- Autonomous-vehicle litigation dossier with a short liability explainer (*updated each offering*)

Week 11 (10/30) — Autonomy and Security

Monday — Governing Autonomous Agents

- Christian (2020), *The Alignment Problem*, prologue and chapter 1 (excerpt)
- De Vynck (2026), “One of Sci-Fi’s Most Difficult Questions about AI Is Becoming Real,” *The Washington Post* (*updated each offering*)

Wednesday — AI as Attack Surface

- CISA and NCSC (2023), *Guidelines for Secure AI System Development*, skim

Part V: Democracy and Development

Week 12 (11/6) — AI and Democratic Institutions

- New York State Office of the Attorney General (2021), *Fake Comments*, executive summary
- Current explainer on AI and election integrity; NCSL deepfake election-law tracker, skim (*updated each offering*)

Week 13 (11/13) — Responsible Development and the Future of Work

Monday — Responsible Development Across Sectors

- Cyberspace Administration of China (2022), “Internet Information Service Algorithmic Recommendation Management Provisions,” DigiChina translation (excerpt)

Wednesday — AI, Work, and Inequality

- Brynjolfsson (2022), “The Turing Trap,” *Daedalus* 151(2)

Part VI: Synthesis

Weeks 14–16 — Audit Fair, Synthesis, Finals

No new reading. Week 14 is the audit fair; Week 15 is synthesis and the final report workshop; Week 16 is finals week.

A Note on How These Readings Were Chosen

The list alternates disciplinary registers on purpose, because this course serves students in public administration and political science alongside students in computer science and cybersecurity.

Public administration and policy anchors come from *Public Administration Review*, the *American Review of Public Administration*, *Perspectives on Public Management and Governance*, and law reviews. Computing anchors come from FAccT, CHI, *Communications of the ACM*, *Science*, the NIST AI Risk Management Framework, and CISA secure-AI guidance.

No reading assumes programming. Technical readings are drawn from the framework and documentation genres — model cards, risk-management frameworks, security guidance — rather than mathematical or systems papers. They are assigned as documentation-literacy exercises that feed directly into the audit project. The one empirical technical paper, Obermeyer et al., is excerpted through its main results, paired with narrative journalism covering the same system, and supported by a reading guide.

Full Citations

Complete bibliographic information for every assigned reading, in alphabetical order. Items described above as updated each offering are posted to *Canvas* at the start of the semester and are not listed here.

Amnesty International. 2021. Xenophobic Machines: Discrimination through Unregulated Use of Algorithms in the Dutch Childcare Benefits Scandal. Technical Report EUR 35/4686/2021 Amnesty International. Dossier B; excerpt assigned.

URL: <https://www.amnesty.org/en/documents/eur35/4686/2021/en/>

Angwin, Julia, Jeff Larson, Surya Mattu and Lauren Kirchner. 2016. "Machine Bias." ProPublica. Dossier A; paired with the Northpointe technical rebuttal exchange.

URL: <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>

Bovens, Mark and Stavros Zouridis. 2002. "From Street-Level to System-Level Bureaucracies: How Information and Communication Technology Is Transforming Administrative Discretion and Constitutional Control." *Public Administration Review* 62(2):174–184.

Bradford, Anu. 2023. *Digital Empires: The Global Battle to Regulate Technology*. New York, NY: Oxford University Press. Three-empires framework chapter assigned as excerpt.

Brynjolfsson, Erik. 2022. "The Turing Trap: The Promise & Peril of Human-Like Artificial Intelligence." *Daedalus* 151(2):272–287.

Bullock, Justin B. 2019. "Artificial Intelligence, Discretion, and Bureaucracy." *The American Review of Public Administration* 49(7):751–761.

Christian, Brian. 2020. *The Alignment Problem: Machine Learning and Human Values*. New York, NY: W. W. Norton. Prologue and chapter 1 assigned as excerpt.

Colorado General Assembly. 2024. "Consumer Protections for Artificial Intelligence (SB24-205)." Colorado Revised Statutes tit. 6, art. 1, pt. 17.

URL: <https://leg.colorado.gov/bills/sb24-205>

Cybersecurity and Infrastructure Security Agency and National Cyber Security Centre. 2023. Guidelines for Secure AI System Development. Technical report CISA and

NCSC. Assigned as skim.

URL: <https://www.cisa.gov/resources-tools/resources/guidelines-secure-ai-system-development>

De Vynck, Gerrit. 2026. “One of Sci-Fi’s Most Difficult Questions about AI Is Becoming Real.” *The Washington Post*.

URL: <https://www.washingtonpost.com/technology/2026/07/13/one-sci-fis-most-difficult-questions-about-ai-is-becoming-real/>

DigiChina. 2022. “Translation: Internet Information Service Algorithmic Recommendation Management Provisions—Effective March 1, 2022.” Stanford University Cyber Policy Center.

URL: <https://digichina.stanford.edu/work/translation-internet-information-service-algorithmic-recommendation-management-provisions-effective-march-1-2022/>

Engstrom, David Freeman, Daniel E. Ho, Catherine M. Sharkey and Mariano-Florentino Cuéllar. 2020. *Government by Algorithm: Artificial Intelligence in Federal Administrative Agencies*. Technical report Administrative Conference of the United States. Executive summary and one agency case study assigned.

URL: <https://www.acus.gov/report/government-algorithm-artificial-intelligence-federal-administrative-agencies>

Eubanks, Virginia. 2018. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. New York, NY: St. Martin’s Press. Required book; purchase or library reserve.

Hogan, Mél. 2015. “Data Flows and Water Woes: The Utah Data Center.” *Big Data & Society* 2(2).

Main, Douglas. 2025. “‘The City That Draws the Line’: One Arizona Community’s Fight against a Huge Datacenter.” *The Guardian*. Tucson “Project Blue” siting dossier.

URL: <https://www.theguardian.com/us-news/2025/oct/15/tucson-arizona-ai-data-center-project-blue>

Mitchell, Margaret, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji and Timnit Gebru. 2019. Model Cards for Model Reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* ’19)*. New York, NY: Association for Computing Machinery pp. 220–229.

Narayanan, Arvind and Sayash Kapoor. 2024. *AI Snake Oil: What Artificial Intelligence Can Do, What It Can't, and How to Tell the Difference*. Princeton, NJ: Princeton University Press. Introduction assigned as excerpt.

National Conference of State Legislatures. N.d. "Artificial Intelligence 2025 Legislation." NCSL.

URL: <https://www.ncsl.org/technology-and-communication/artificial-intelligence-2025-legislation>

National Institute of Standards and Technology. 2023. Artificial Intelligence Risk Management Framework (AI RMF 1.0). Technical Report NIST AI 100-1 U.S. Department of Commerce. Core functions assigned as skim.

URL: <https://www.nist.gov/itl/ai-risk-management-framework>

New York State Office of the Attorney General. 2021. Fake Comments: How U.S. Companies & Partisans Hack Democracy to Undermine Your Voice. Technical report Office of the New York State Attorney General. Executive summary assigned.

URL: <https://ag.ny.gov/sites/default/files/oag-fakecomments-report.pdf>

Nicoletti, Leonardo, Michelle Ma and Dina Bass. 2025. "The AI Boom Is Draining Water from the Areas That Need It Most." Bloomberg Graphics. Finds roughly two-thirds of US data centers built or in development since 2022 sit in high-water-stress areas.

URL: <https://www.bloomberg.com/graphics/2025-ai-impacts-data-centers-water-data/>

Obermeyer, Ziad, Brian Powers, Christine Vogeli and Sendhil Mullainathan. 2019. "Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations." *Science* 366(6464):447–453. Assigned through main results; methods appendix optional.

Pahlka, Jennifer. 2023. *Recoding America: Why Government Is Failing in the Digital Age and How We Can Do Better*. New York, NY: Metropolitan Books. One chapter assigned as excerpt.

Ross, Casey and Bob Herman. 2023. "Denied by AI: How Medicare Advantage Plans Use Algorithms to Cut Off Care for Seniors in Need." STAT News. First installment of the four-part "Denied by AI" series; Pulitzer Prize finalist.

URL: <https://www.statnews.com/2023/03/13/medicare-advantage-plans-denial-artificial-intelligence/>

Veale, Michael and Frederik Zuiderveen Borgesius. 2021. "Demystifying the Draft EU

Artificial Intelligence Act: Analysing the Good, the Bad, and the Unclear Elements of the Proposed Approach.” *Computer Law Review International* 22(4):97–112.

Wheeler, Tom. 2023. “The Three Challenges of AI Regulation.” Brookings Institution.

URL: <https://www.brookings.edu/articles/the-three-challenges-of-ai-regulation/>